



## Problem Statement : Build CNN Model for Classification Of Flowers

```
!gdown --id 1xkynpl15pt6KT3YS1Dimu4A5iRU9qYck
```

```
/usr/local/lib/python3.7/dist-packages/gdown/cli.py:131: FutureWarning: Option `--id`  
category=FutureWarning,  
Downloading...  
From: https://drive.google.com/uc?id=1xkynpl15pt6KT3YS1Dimu4A5iRU9qYck  
To: /content/Flowers-Dataset.zip  
100% 236M/236M [00:00<00:00, 285MB/s]
```

```
!unzip Flowers-Dataset.zip
```

```
inflating: flowers/tulip/8712270243_8512cf4fbd.jpg  
inflating: flowers/tulip/8712270665_57b5bda0a2_n.jpg  
inflating: flowers/tulip/8712282563_3819afb7bc.jpg  
inflating: flowers/tulip/8713357842_9964a93473_n.jpg  
inflating: flowers/tulip/8713387500_6a9138b41b_n.jpg  
inflating: flowers/tulip/8713388322_e5ae26263b_n.jpg  
inflating: flowers/tulip/8713389178_66bceb71a8_n.jpg  
inflating: flowers/tulip/8713390684_041148dd3e_n.jpg  
inflating: flowers/tulip/8713391394_4b679ea1e3_n.jpg  
inflating: flowers/tulip/8713392604_90631fb809_n.jpg  
inflating: flowers/tulip/8713394070_b24561b0a9.jpg  
inflating: flowers/tulip/8713396140_5af8136136.jpg  
inflating: flowers/tulip/8713397358_0505cc0176_n.jpg  
inflating: flowers/tulip/8713397694_bcbcbba2c2_n.jpg  
inflating: flowers/tulip/8713398114_bc96f1b624_n.jpg  
inflating: flowers/tulip/8713398614_88202e452e_n.jpg  
inflating: flowers/tulip/8713398906_28e59a225a_n.jpg  
inflating: flowers/tulip/8713407768_f880df361f.jpg  
inflating: flowers/tulip/8717900362_2aa508e9e5.jpg  
inflating: flowers/tulip/8722514702_7ecc68691c.jpg  
inflating: flowers/tulip/8723767533_9145dec4bd_n.jpg  
inflating: flowers/tulip/8729501081_b993185542_m.jpg  
inflating: flowers/tulip/8733586143_3139db6e9e_n.jpg  
inflating: flowers/tulip/8748266132_5298a91dcf_n.jpg  
inflating: flowers/tulip/8750288831_5e49a9f29b.jpg  
inflating: flowers/tulip/8757486380_90952c5377.jpg  
inflating: flowers/tulip/8758464923_75a5ffe320_n.jpg  
inflating: flowers/tulip/8758519201_16e8d2d781_n.jpg  
inflating: flowers/tulip/8759594528_2534c0ec65_n.jpg  
inflating: flowers/tulip/8759597778_7fca5d434b_n.jpg  
inflating: flowers/tulip/8759601388_36e2a50d98_n.jpg  
inflating: flowers/tulip/8759606166_8e475013fa_n.jpg  
inflating: flowers/tulip/8759618746_f5e39fdbf8_n.jpg  
inflating: flowers/tulip/8762189906_8223cef62f.jpg  
inflating: flowers/tulip/8762193202_0fbf2f6a81.jpg  
inflating: flowers/tulip/8768645961_8f1e097170_n.jpg
```

```

inflating: flowers/tulip/8817622133_a42bb90e38_n.jpg
inflating: flowers/tulip/8838347159_746d14e6c1_m.jpg
inflating: flowers/tulip/8838354855_c474fc66a3_m.jpg
inflating: flowers/tulip/8838914676_8ef4db7f50_n.jpg
inflating: flowers/tulip/8838975946_f54194894e_m.jpg
inflating: flowers/tulip/8838983024_5c1a767878_n.jpg
inflating: flowers/tulip/8892851067_79242a7362_n.jpg
inflating: flowers/tulip/8904780994_8867d64155_n.jpg
inflating: flowers/tulip/8908062479_449200a1b4.jpg
inflating: flowers/tulip/8908097235_c3e746d36e_n.jpg
inflating: flowers/tulip/9019694597_2d3bbbedb17.jpg
inflating: flowers/tulip/9030467406_05e93ff171_n.jpg
inflating: flowers/tulip/9048307967_40a164a459_m.jpg
inflating: flowers/tulip/924782410_94ed7913ca_m.jpg
inflating: flowers/tulip/9378657435_89fabf13c9_n.jpg
inflating: flowers/tulip/9444202147_405290415b_n.jpg
inflating: flowers/tulip/9446982168_06c4d71da3_n.jpg
inflating: flowers/tulip/9831362123_5aac525a99_n.jpg
inflating: flowers/tulip/9870557734_88eb3b9e3b_n.jpg
inflating: flowers/tulip/9947374414_fdf1d0861c_n.jpg
inflating: flowers/tulip/9947385346_3a8cacea02_n.jpg
inflating: flowers/tulip/9976515506_d496c5e72c.jpg

```

```

import numpy as np
import pandas as pd
import os
import torch
import torchvision
import tarfile
import torchvision
from torch.utils.data import random_split
from torchvision.datasets import ImageFolder
from torchvision import transforms
from torchvision.transforms import ToTensor
from torch.utils.data.dataloader import DataLoader
import torch.nn as nn
from torchvision.utils import make_grid
import torchvision.models as models
import torch.nn.functional as F
import matplotlib.pyplot as plt
%matplotlib inline

data_dir ="flowers"

```

## ▼ Applying Data Augmentation

```

transformer = torchvision.transforms.Compose(
    [
        torchvision.transforms.Resize((224, 224)),
        torchvision.transforms.RandomHorizontalFlip(p=0.5),
        torchvision.transforms.RandomVerticalFlip(p=0.5),
        torchvision.transforms.RandomRotation(40),
    ]
)

```

```

        torchvision.transforms.ToTensor(),
        torchvision.transforms.Normalize(
            mean=[0.4914, 0.4822, 0.4465], std=[0.2023, 0.1994, 0.2010]
        ),
    ]
)
database = ImageFolder(data_dir, transform=transformer)

database.classes

['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']

flower_name_dict={'daisy':0 ,
                  'dandelion':1,
                  'rose':2 ,
                  'sunflower':3 ,
                  'tulip':4}

def encode_label(img_label):
    return flower_name_dict[img_label]

flower_label = {
    0: 'daisy',
    1: 'dandelion',
    2: 'rose',
    3: 'sunflower',
    4: 'tulip'
}


def show_batch(dl,invert=True):
    for images, labels in dl:
        fig, ax = plt.subplots(figsize=(12, 6))
        ax.set_xticks([]); ax.set_yticks([])
        ax.imshow(make_grid(images, nrow=16).permute(1, 2, 0))
        break

def show_sample(image, label,invert=True):
    print("Label :" +database.classes[label] + "(" + str(label) + ")")
    plt.imshow(image.permute(1, 2, 0))

show_sample(*database[1])

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB c  
Label :daisy(0)



## ▼ Splitting the data for validation dataset and train dataset

---

```
validation_size = 500
training_size = len(database) - validation_size
```

0 50 100 150 200

```
train_ds, val_ds_main = random_split(database,[training_size, validation_size])
val_ds, test_ds = random_split(val_ds_main,[300, 200])
len(train_ds), len(val_ds)
```


(3817, 300)

batch\_size=32

```
train_dl = DataLoader(train_ds, batch_size , shuffle=True)
val_dl = DataLoader(val_ds,batch_size)
test_dl = DataLoader(test_ds, batch_size)
```

show\_batch(train\_dl,invert=True)

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB c



## ▼ 1. Training the model for image classification

```
def accuracy(outputs, labels):
    _, preds = torch.max(outputs, dim=1)
    return torch.tensor(torch.sum(preds == labels).item() / len(preds))
```

```
class ImageClassification(nn.Module):
    def training_step(self, batch):
        images, labels = batch
        out = self(images) # Generate predictions
        loss = F.cross_entropy(out, labels) # Calculate loss
        return loss
```

```

def validation_step(self, batch):
    images, labels = batch
    out = self(images)                # Generate predictions
    loss = F.cross_entropy(out, labels) # Calculate loss
    acc = accuracy(out, labels)        # Calculate accuracy
    return {'val_loss': loss.detach(), 'val_acc': acc}

def validation_epoch_end(self, outputs):
    batch_losses = [x['val_loss'] for x in outputs]
    epoch_loss = torch.stack(batch_losses).mean() # Combine losses
    batch_accs = [x['val_acc'] for x in outputs]
    epoch_acc = torch.stack(batch_accs).mean()    # Combine accuracies
    return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}

def epoch_end(self, epoch, result):
    print("Epoch [{}], train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(
        epoch, result['train_loss'], result['val_loss'], result['val_acc']))

class FlowerModel(ImageClassification):
    def __init__(self):
        super().__init__()
        self.network = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.Conv2d(256, 256, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Flatten(),
            nn.Linear(256*28*28, 1024),
            nn.ReLU(),
            nn.Linear(1024, 512),
            nn.ReLU(),
            nn.Linear(512, 5))

    def forward(self, xb):
        return self.network(xb)

model = FlowerModel()
model

```

```

FlowerModel(
  (network): Sequential(
    (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU()
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU()
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU()
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU()
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (15): Flatten(start_dim=1, end_dim=-1)
    (16): Linear(in_features=200704, out_features=1024, bias=True)
    (17): ReLU()
    (18): Linear(in_features=1024, out_features=512, bias=True)
    (19): ReLU()
    (20): Linear(in_features=512, out_features=5, bias=True)
  )
)

```

```

for images, labels in train_dl:
    print('images.shape:', images.shape)
    out = model(images)
    print('out.shape:', out.shape)
    print('out[0]:', out[0])
    break

```

```

images.shape: torch.Size([32, 3, 224, 224])
out.shape: torch.Size([32, 5])
out[0]: tensor([ 0.0472,  0.0452, -0.0277, -0.0230,  0.0314],
               grad_fn=<SelectBackward0>)

```

```

@torch.no_grad()
def evaluate(model, val_loader):
    model.eval()
    outputs = [model.validation_step(batch) for batch in val_loader]
    return model.validation_epoch_end(outputs)

def fit(epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.SGD):
    history = []
    optimizer = opt_func(model.parameters(), lr)
    for epoch in range(epochs):
        # Training Phase
        model.train()
        train_losses = []
        for batch in train_loader:
            loss = model.training_step(batch)
            train_losses.append(loss)
            loss.backward()

```

```

        optimizer.step()
        optimizer.zero_grad()
    # Validation phase
    result = evaluate(model, val_loader)
    result['train_loss'] = torch.stack(train_losses).mean().item()
    model.epoch_end(epoch, result)
    history.append(result)
return history

```

## ▼ Get GPU up on running

```

def get_default_device():
    """Pick GPU if available, else CPU"""
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    """Move tensor(s) to chosen device"""
    if isinstance(data, (list,tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    """Wrap a dataloader to move data to a device"""
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        """Yield a batch of data after moving it to device"""
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        """Number of batches"""
        return len(self.dl)

device = get_default_device()

device = get_default_device()
device

device(type='cuda')

```

## ▼ Training and Validation Datasets

```
train_dl = DeviceDataLoader(train_dl, device)
val_dl = DeviceDataLoader(val_dl, device)
to_device(model, device);

model = to_device(FlowerModel(), device)

evaluate(model, val_dl)

{'val_loss': 1.6087734699249268, 'val_acc': 0.23854167759418488}

num_epochs = 10
opt_func = torch.optim.Adam
lr = 0.001


%%time
try1 = fit(num_epochs, lr, model, train_dl, val_dl, opt_func)

Epoch [0], train_loss: 1.5009, val_loss: 1.3029, val_acc: 0.4271
Epoch [1], train_loss: 1.3080, val_loss: 1.3151, val_acc: 0.4365
Epoch [2], train_loss: 1.2726, val_loss: 1.2642, val_acc: 0.4333
Epoch [3], train_loss: 1.2631, val_loss: 1.2681, val_acc: 0.4625
Epoch [4], train_loss: 1.2557, val_loss: 1.2266, val_acc: 0.4594
Epoch [5], train_loss: 1.1785, val_loss: 1.1668, val_acc: 0.5406
Epoch [6], train_loss: 1.0677, val_loss: 1.1469, val_acc: 0.5479
Epoch [7], train_loss: 1.0231, val_loss: 1.0959, val_acc: 0.5542
Epoch [8], train_loss: 0.9870, val_loss: 1.0238, val_acc: 0.5802
Epoch [9], train_loss: 0.9820, val_loss: 1.0775, val_acc: 0.6104
CPU times: user 6min 57s, sys: 4.41 s, total: 7min 2s
Wall time: 7min 2s

def plot_accuracies(try1):
    accuracies = [x['val_acc'] for x in try1]
    plt.plot(accuracies, '-x')
    plt.xlabel('epoch')
    plt.ylabel('accuracy')
    plt.title('Accuracy vs. No. of epochs');

plot_accuracies(try1)
```






```
def plot_losses(try1):
    train_losses = [x.get('train_loss') for x in try1]
    val_losses = [x['val_loss'] for x in try1]
    plt.plot(train_losses, '-bx')
    plt.plot(val_losses, '-rx')
    plt.xlabel('epoch')
    plt.ylabel('loss')
    plt.legend(['Training', 'Validation'])
    plt.title('Loss vs. No. of epochs');

    epoch
```

```
plot_losses(try1)
```



```
test_dl = DeviceDataLoader(test_dl, device)
evaluate(model, test_dl)
```


```
{'val_loss': 0.9322658777236938, 'val_acc': 0.5982142686843872}
```

## ▼ Predict the image

```
def predict_image(img, model):
    xb = img.unsqueeze(0)
    yb = model(xb)
    _, preds = torch.max(yb, dim=1)
    return flower_label[preds[0].item()]
```

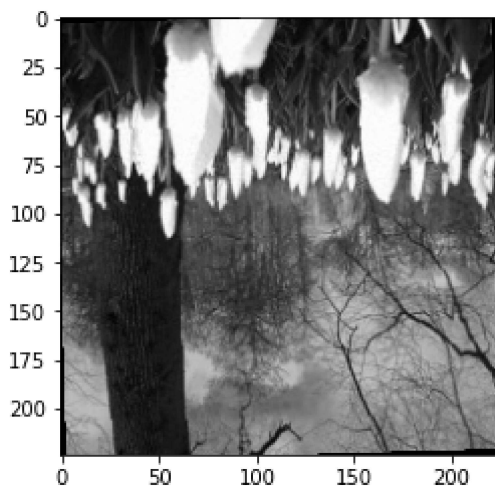
```
img, label = test_ds[1]
plt.imshow(img[0], cmap='gray')
print('Label:', database.classes[label], ', Predicted:', predict_image(img, FlowerModel()))
```

Label: dandelion , Predicted: dandelion



```
img, label = test_ds[5]
plt.imshow(img[0], cmap='gray')
print('Label:', database.classes[label], ', Predicted:', predict_image(img, FlowerModel()))
```

Label: tulip , Predicted: sunflower



## ▼ Saving the model

```
torch.save(model.state_dict(), 'Flower-classification.pt')
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 3s completed at 9:48 AM

