

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

Loading the dataset provided in assignment folder

```
df=pd.read_csv("Churn_Modelling.csv")
df.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore             10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                    10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                 10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard              10000 non-null  int64
11  IsActiveMember         10000 non-null  int64
12  EstimatedSalary        10000 non-null  float64
13  Exited                 10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```


df.describe()

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.00000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.914119
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.285622
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000

1. UNIVARIATE ANALYSIS


Histogram






```
sns.histplot(y="Age",data=df,color='darkorange')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8c97b6fd10>



```
sns.histplot(x='Age',data=df,hue=df['Tenure'])
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f8c975fbfd0>



Distplot


```
sns.distplot(df["Age"],color='violet')
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f8c975b4450>



```
sns.distplot(df["Age"],hist=False,color='darkblue')
```

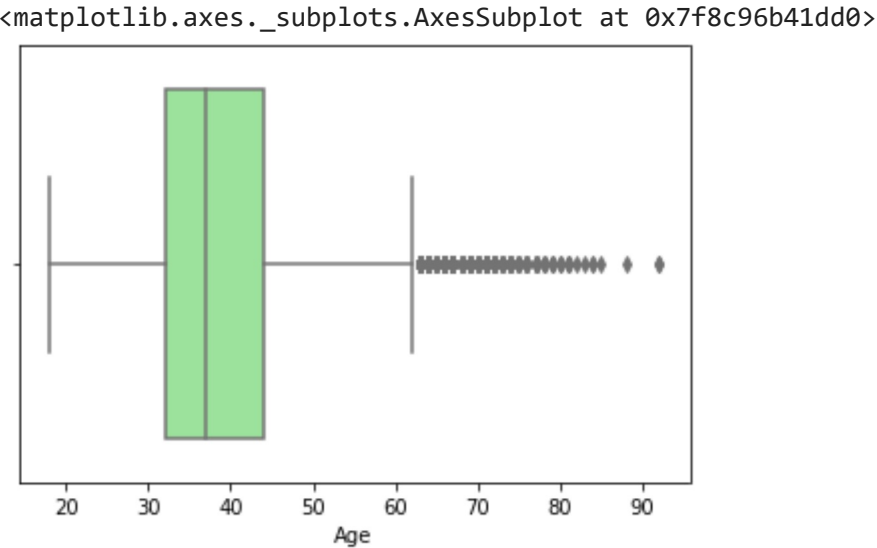
<matplotlib.axes._subplots.AxesSubplot at 0x7f8c97624990>






Boxplot

```
sns.boxplot(df["Age"],color='lightgreen')
```



Countplot

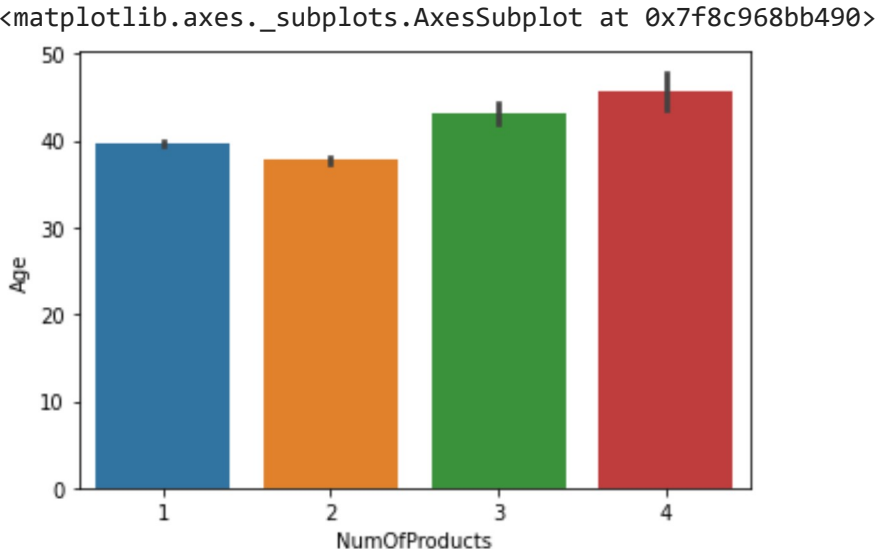
```
sns.countplot(df['Age'])
```



2. BIVARIATE ANALYSIS

Barplot


```
sns.barplot(df["NumOfProducts"],df["Age"])
```



Linearplot


```
sns.lineplot(df["Age"],df["NumOfProducts"], color='purple')
```





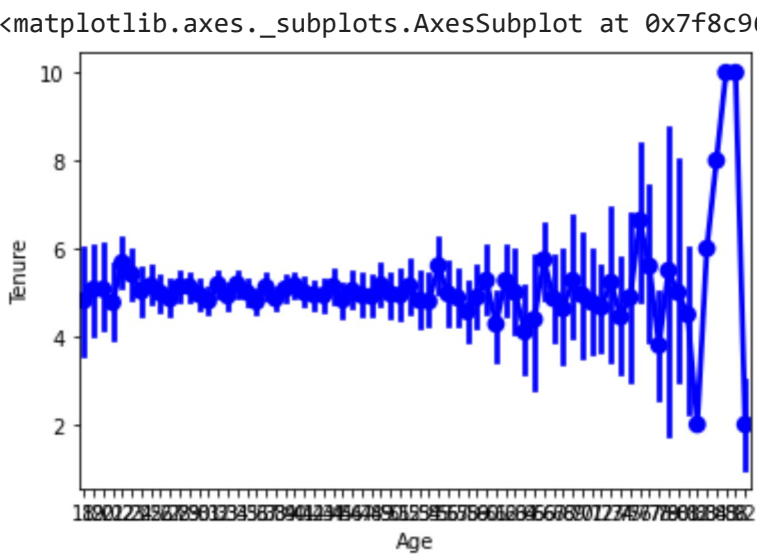
Scatterplot

```
sns.scatterplot(x=df.Age,y=df.RowNumber,color='pink')
```



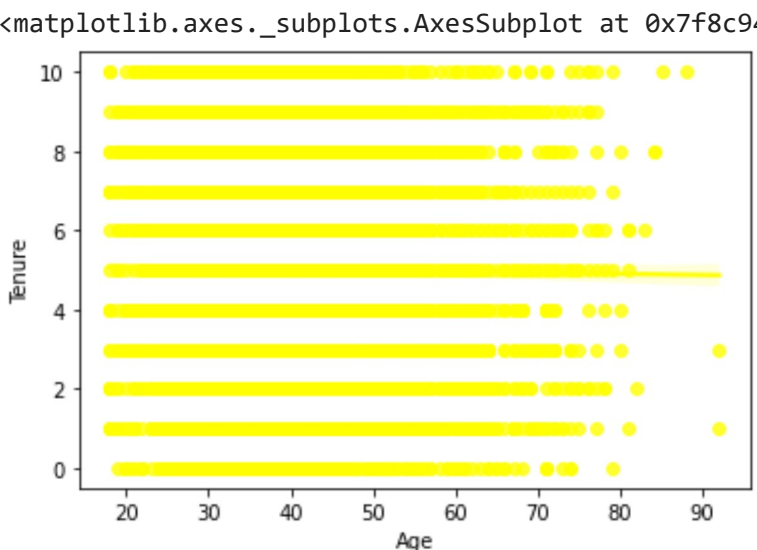
Pointplot

```
sns.pointplot(x='Age',y='Tenure',data=df,color='blue')
```



Regplot

```
sns.regplot(df['Age'],df['Tenure'],color='yellow')
```

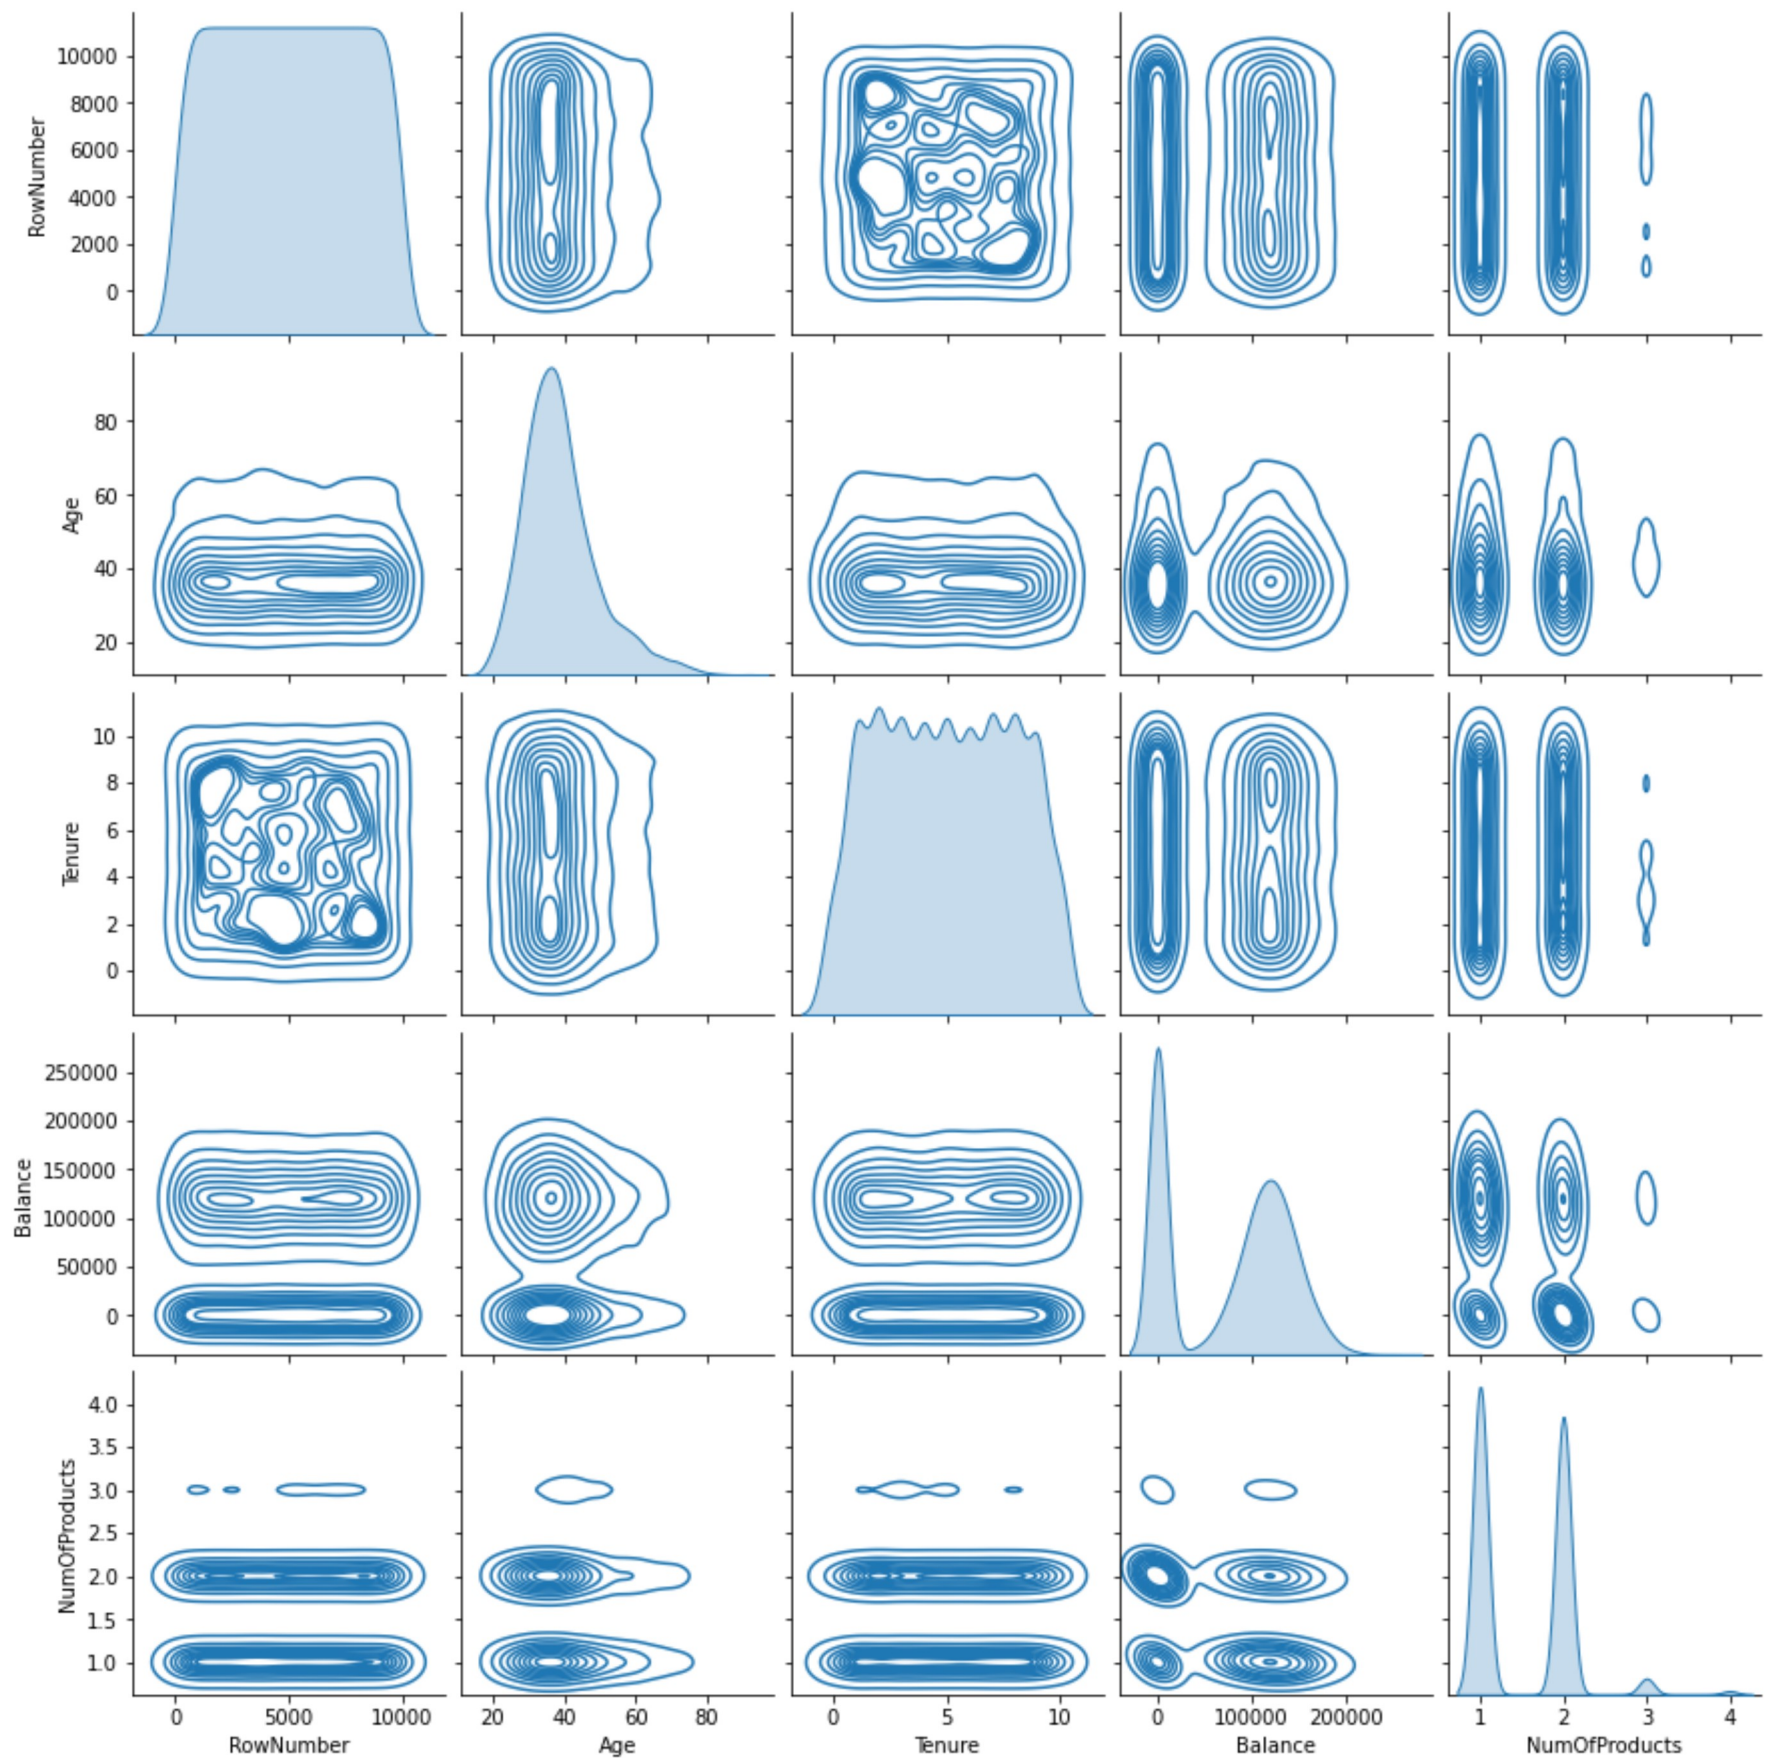


3. MULTI - VARIATE ANALYSIS

Pairplot

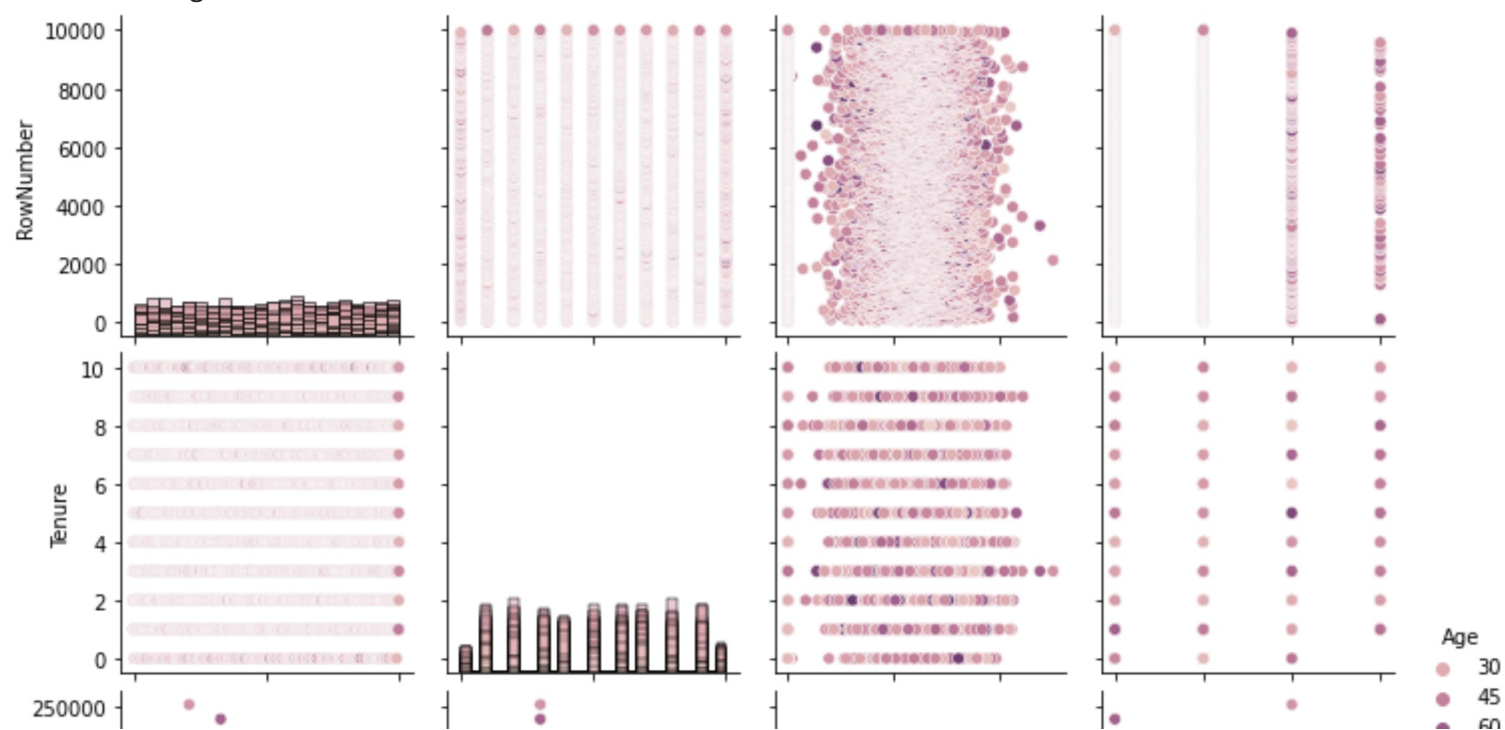
```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], kind="kde")
```

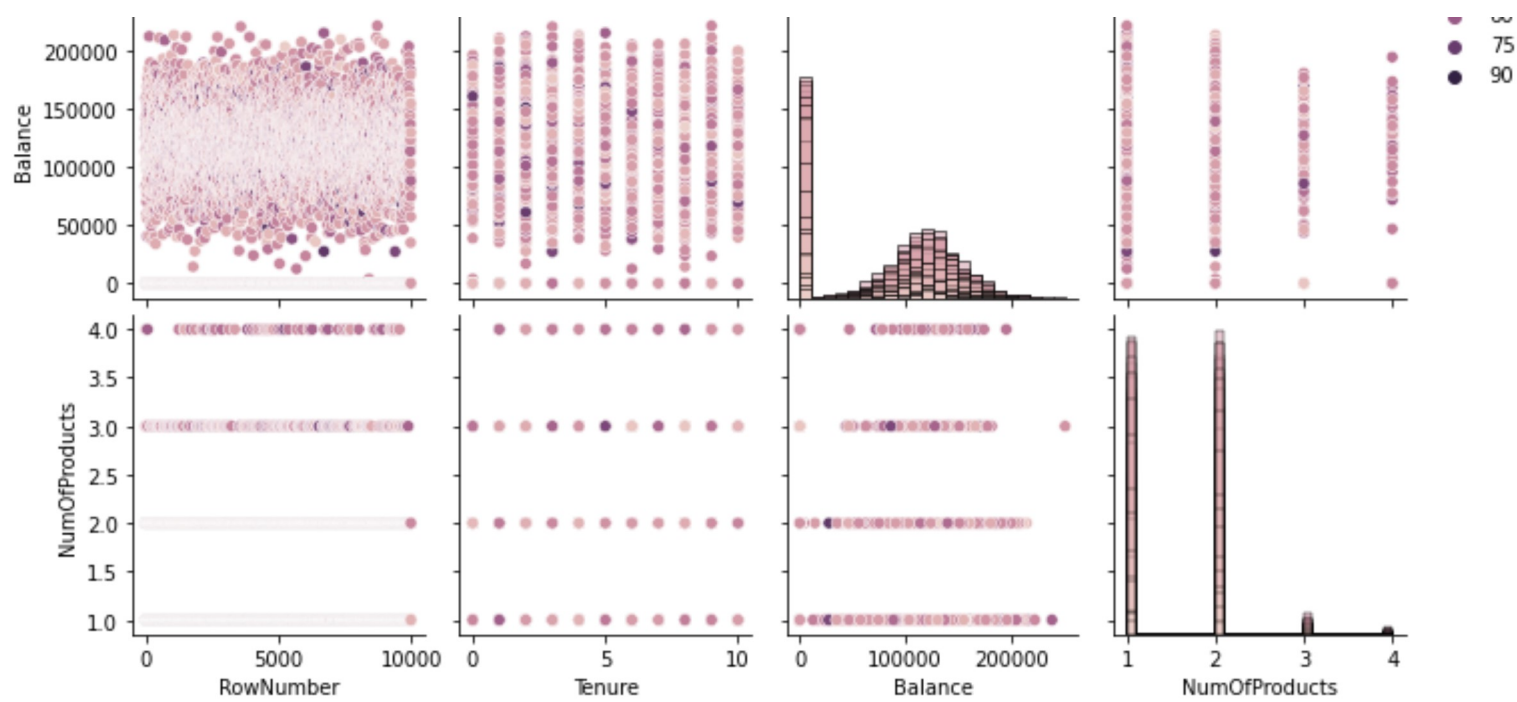
<seaborn.axisgrid.PairGrid at 0x7f8c94cb01d0>



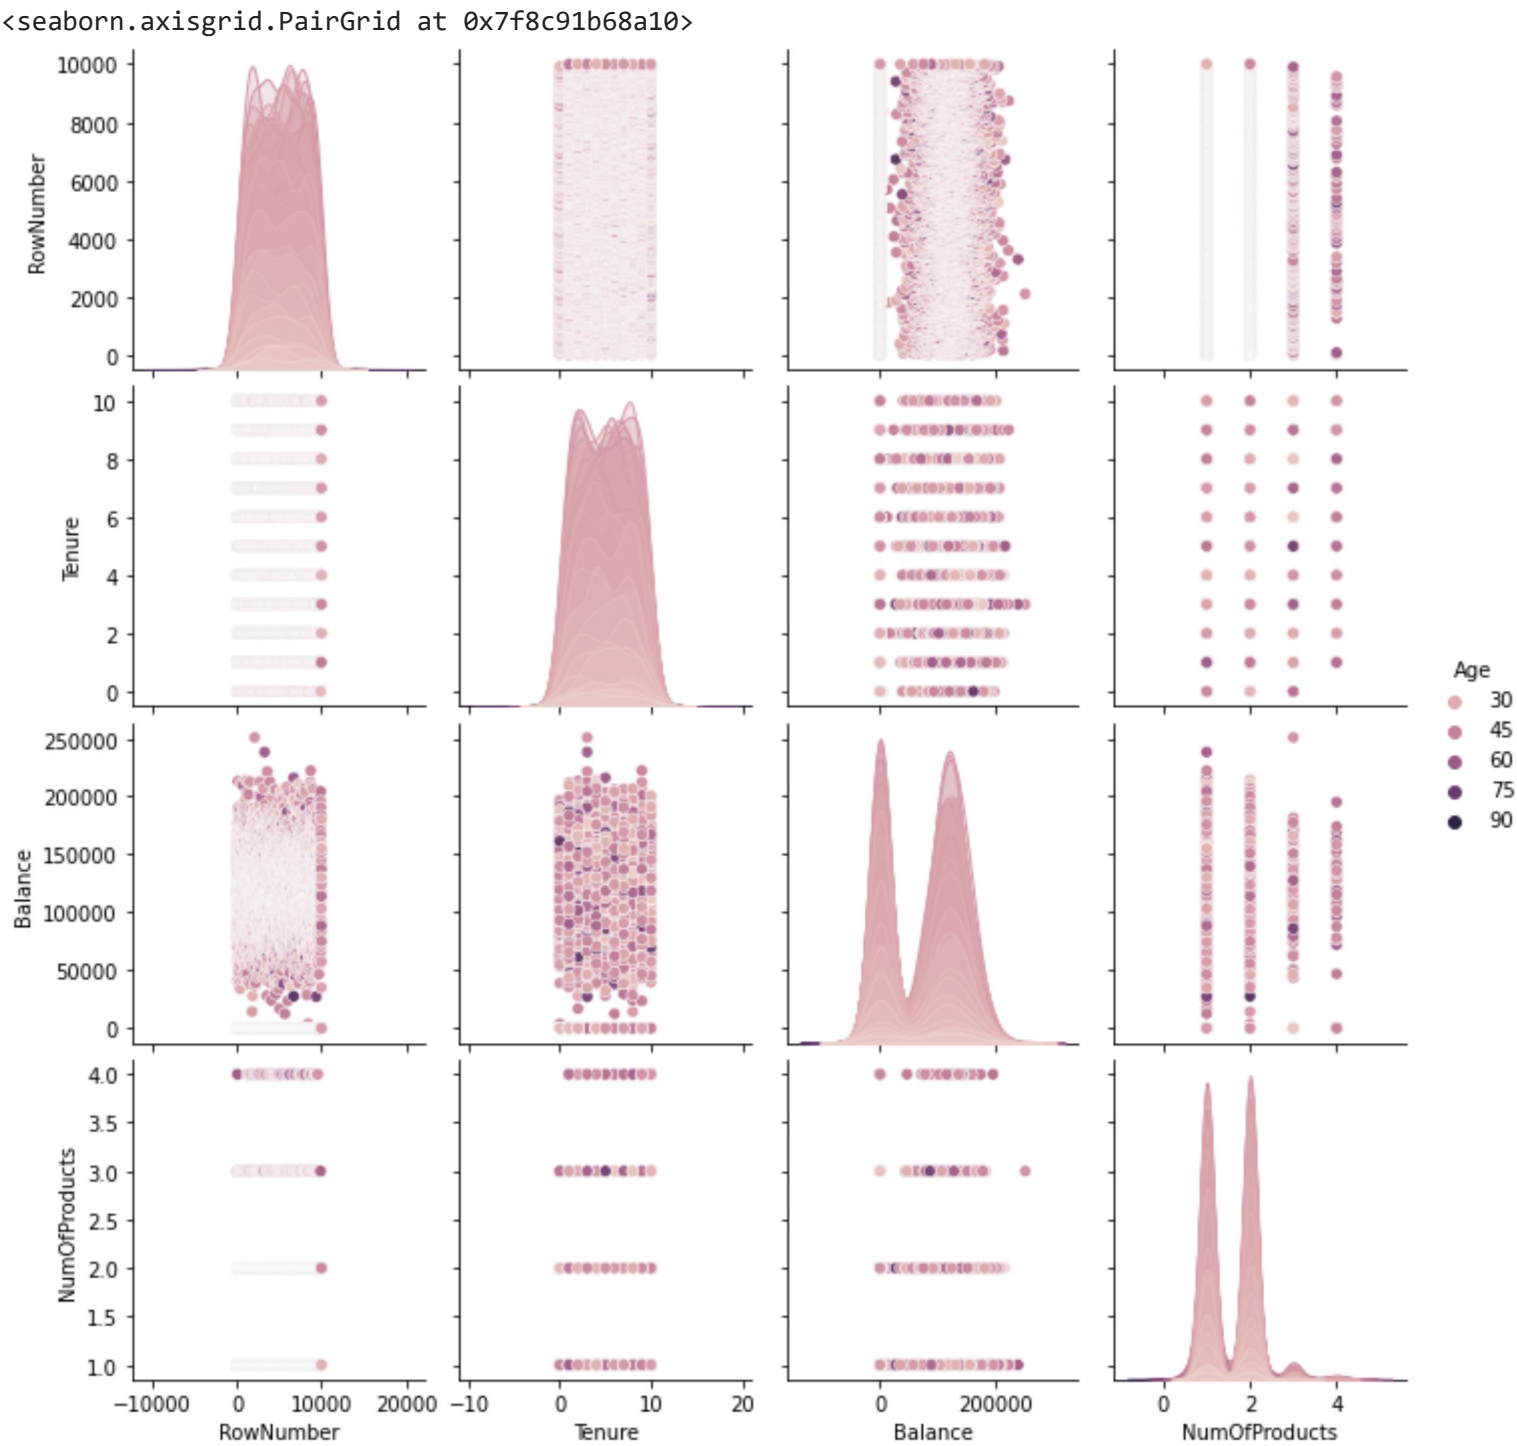
```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], hue="Age", diag_kind="hist")
```

<seaborn.axisgrid.PairGrid at 0x7f8c93e736d0>





```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], hue="Age")
```



4. Perform descriptive statistics on the dataset

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.919200
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.420600
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000

75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000



5. Handle the Missing values.

```
data=pd.DataFrame({"a":[1,2,np.nan], "b":[1,np.nan,np.nan], "c":[1,2,4]})
data
```

	a	b	c
0	1.0	1.0	1
1	2.0	NaN	2
2	NaN	NaN	4



```
data.isnull().any()
```

a True
b True
c False
dtype: bool

```
data.isnull().sum()
```

a 1
b 2
c 0
dtype: int64

```
data.fillna(value = "S")
```

	a	b	c
0	1.0	1.0	1
1	2.0	S	2
2	S	S	4



```
data["a"].mean()
```

1.5

```
data["a"].median()
```

1.5

6. Find the outliers and replace the outliers

```
outliers=df.quantile(q=(0.25,0.75))
```

```
outliers
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0.25	2500.75	15628528.25	584.0	32.0	3.0	0.00	1.0	0.0	0.0	51002.1100
0.75	7500.25	15753233.75	718.0	44.0	7.0	127644.24	2.0	1.0	1.0	149388.2475

```
aaa=outliers.loc[0.75]-outliers.loc[0.25]
```

```
aaa
```

RowNumber 4999.5000
CustomerId 124705.5000
CreditScore 134.0000
Age 12.0000


```
-----
Tenure      4.0000
Balance    127644.2400
NumOfProducts  1.0000
HasCrCard   1.0000
IsActiveMember  1.0000
EstimatedSalary  98386.1375
Exited      0.0000
dtype: float64
```


```
low = outliers.loc[0.25] - 1.5*aaa
```

```
low
RowNumber      -4.998500e+03
CustomerId     1.544147e+07
CreditScore    3.830000e+02
Age            1.400000e+01
Tenure         -3.000000e+00
Balance        -1.914664e+05
NumOfProducts  -5.000000e-01
HasCrCard      -1.500000e+00
IsActiveMember -1.500000e+00
EstimatedSalary -9.657710e+04
Exited         0.000000e+00
dtype: float64
```

```
high = outliers.loc[0.75] + 1.5 * aaa
high
```

```
RowNumber      1.499950e+04
CustomerId     1.594029e+07
CreditScore    9.190000e+02
Age            6.200000e+01
Tenure         1.300000e+01
Balance        3.191106e+05
NumOfProducts  3.500000e+00
HasCrCard      2.500000e+00
IsActiveMember 2.500000e+00
EstimatedSalary 2.969675e+05
Exited         0.000000e+00
dtype: float64
```


```
sns.boxplot(df['Age']);
```



```
df["Age"]=np.where(df["Age"]<25,50,df["Age"])
```

```
sns.boxplot(df["Age"],color='pink')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8c8d6033d0>
```



7. Check for Categorical columns and perform encoding.

```
df.head(4)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	

```
df["Gender"].replace({"Female":0,"Male":1},inplace = True)
df["Geography"].replace({"France":1,"Spain":2,"Germany":3},inplace = True)
df["Gender"].replace({"Female":0,"Male":1},inplace = True)
df["Geography"].replace({"France":1,"Spain":2,"Germany":3},inplace = True)
```

```
df.head(4)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15634602	Hargrave	619	1	0	42	2	0.00	1	1	
1	2	15647311	Hill	608	2	0	41	1	83807.86	1	0	
2	3	15619304	Onio	502	1	0	42	8	159660.80	3	1	
3	4	15701354	Boni	699	1	0	39	1	0.00	2	0	



8. Split the data into dependent and independent variables.

```
y = df["Surname"]
x=df.drop(columns=["Surname"],axis=1)
x.head()
```

	RowNumber	CustomerId	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	619	1	0	42	2	0.00	1	1	1	40461.56
1	2	15647311	608	2	0	41	1	83807.86	1	0	1	59190.62
2	3	15619304	502	1	0	42	8	159660.80	3	1	0	113547.53
3	4	15701354	699	1	0	39	1	0.00	2	0	0	40461.56
4	5	15737888	850	2	0	43	2	125510.82	1	1	1	113547.53

9. Scale the independent variables

```
names=x.columns
names

Index(['RowNumber', 'CustomerId', 'CreditScore', 'Geography', 'Gender', 'Age',
      'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember',
      'EstimatedSalary', 'Exited'],
      dtype='object')
```

```
from sklearn.preprocessing import scale
X=scale(x)
X

array([[ -1.73187761, -0.78321342, -0.32622142, ...,  0.97024255,
         0.02188649,  1.97716468],
       [ -1.7315312 , -0.60653412, -0.44003595, ...,  0.97024255,
         0.21653375, -0.50577476],
       [ -1.73118479, -0.99588476, -1.53679418, ..., -1.03067011,
         0.2406869 ,  1.97716468],
       ...,
       [  1.73118479,  -1.17028179,  0.60108839, ...,  0.97024255,
```

```
[ 1.73110473, -1.47328173,  0.00430833, ...,  0.37024233,
-1.00864308,  1.97716468],
[ 1.7315312 , -0.11935577,  1.25683526, ..., -1.03067011,
-0.12523071,  1.97716468],
[ 1.73187761, -0.87055909,  1.46377078, ..., -1.03067011,
-1.07636976, -0.50577476]])
```

```
x = pd.DataFrame(X,columns = names )
x
```

	RowNumber	CustomerId	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	-1.731878	-0.783213	-0.326221	-0.902587	-1.095988	0.179622	-1.041760	-1.225848	-0.911583	0.646092	
1	-1.731531	-0.606534	-0.440036	0.301665	-1.095988	0.080092	-1.387538	0.117350	-0.911583	-1.547768	
2	-1.731185	-0.995885	-1.536794	-0.902587	-1.095988	0.179622	1.032908	1.333053	2.527057	0.646092	
3	-1.730838	0.144767	0.501521	-0.902587	-1.095988	-0.118968	-1.387538	-1.225848	0.807737	-1.547768	
4	-1.730492	0.652659	2.063884	0.301665	-1.095988	0.279152	-1.041760	0.785728	-0.911583	0.646092	
...	
9995	1.730492	-1.177652	1.246488	-0.902587	0.912419	-0.118968	-0.004426	-1.225848	0.807737	0.646092	
9996	1.730838	-1.682806	-1.391939	-0.902587	0.912419	-0.517088	1.724464	-0.306379	-0.911583	0.646092	
9997	1.731185	-1.479282	0.604988	-0.902587	-1.095988	-0.417558	0.687130	-1.225848	-0.911583	-1.547768	
9998	1.731531	-0.119356	1.256835	1.505917	0.912419	0.179622	-0.695982	-0.022608	0.807737	0.646092	
9999	1.731878	-0.870559	1.463771	-0.902587	-1.095988	-1.213798	-0.350204	0.859965	-0.911583	0.646092	

10000 rows × 13 columns

10. Split the data into training and testing

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
x_train.head()
```

	RowNumber	CustomerId	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
7389	0.827747	-0.195066	0.170424	0.301665	-1.095988	-0.616618	-0.004426	-1.225848	0.807737	0.646092	
9275	1.481077	0.810821	-2.312802	1.505917	0.912419	0.179622	-1.387538	-0.012892	-0.911583	0.646092	
2995	-0.694379	-1.507642	-1.195351	-0.902587	-1.095988	-1.114268	-1.041760	0.575076	-0.911583	0.646092	
5316	0.109639	1.243462	0.035916	0.301665	0.912419	-0.019438	-0.004426	0.467955	-0.911583	0.646092	
356	-1.608556	-1.100775	2.063884	0.301665	-1.095988	1.672571	1.032908	0.806010	0.807737	0.646092	

```
x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

((8000, 13), (8000,), (2000, 13), (2000,))

